

Impact of Root DNS Latency on User Experience

John Heidemann (ISI)

Matt Calder (Microsoft Research)

Arpit Gupta (UCSB)

Ethan Katz-Bassett (Columbia University)

Thomas Koch (Columbia University)

ABSTRACT

Anycast is means of distributing content that has been praised for its simplicity and performance, yet criticized for inflating client latencies in some cases. The root DNS servers are frequent sources of information for studies analyzing anycast, since the information is relatively easy to obtain. We argue that the root DNS servers are not valid test subjects for studies either offering criticism of or suggesting improvements for anycast, since clients rarely interact with the root DNS infrastructure. We demonstrate that simple caching policies of recursive resolvers limit a clients’ exposure to the root DNS infrastructure, and quantify how much latency a client experiences each day due to root DNS resolution. These results indicate that future studies should not draw from root DNS server data to support arguments regarding anycast latency inflation.

1 INTRODUCTION

Anycast is a method of server deployment in which a set of physically distinct servers, called anycast replicas, all serve the same content. Specifically, IP anycast is a system in which geographically diverse anycast replicas all advertise the same IP address using the Border Gateway Protocol (BGP). Routing policies of providers and BGP determine which anycast replica a client will be mapped to when querying the (anycasted) IP. This is all transparent to the client, as each site has exactly the same content.

Anycast has been lauded for its potential to offer improved latency and decreased load on each anycast server, doing so with minimal scaling complexity. The geographical diversity of anycast deployments offers resilience to network outages and targeted attacks [12]. However, numerous studies have argued that anycast can lead to reduced performance for a subset of clients. Some go further to say that adding more sites may harm, rather than help user performance. The basic idea behind these claims is that BGP can sometimes direct clients to far-away anycast servers, despite the existence of a geographically close server. This inefficiency unnecessarily inflates

latencies for users, and harms their experience. The root DNS servers notoriously feature in studies involving anycast because it is relatively easy to gain access to root DNS data, information about their deployments, and the fact that they are run by several organizations. This last fact manifests itself in a diverse set of deployment strategies, for essentially the same service. For example, in 2018 B-root had 2 server deployments while K-root had close to 70. These different strategies allow researchers to, for example, coarsely analyze the effect the number of sites has on the performance of the service. Using the root DNS servers to draw conclusions about the performance of anycast as a whole has its drawbacks, as several authors note. Organizations in charge of managing these services have different performance goals when deploying new servers. For example, generally, new root server deployments offer more resilience in the face of attacks on the DNS infrastructure and so organizations may opt to place their servers in locations that maximize reachability to at least one replica in the face of server failure. Conversely, managers of large content delivery networks (CDN’s) may opt to place their servers near large population centers to provide low latency access to users. This is on top of the infrastructure the CDN may already have in place, such as a diverse set of peers or a private wide area network (WAN). Indeed the everyday notions of “performance” that typical users are interested in are principally important to CDN’s, yet only somewhat important to root DNS managers. Nevertheless, in the face of perceived anycast inefficiencies many authors have offered solutions ranging from ISP policy changes to modifications to BGP to bolster the performance of anycast. Often, studies have framed the problem as a sort of optimization problem, where every client request should be mapped to the same physical site as the best unicast alternative. These proposed modifications are then shown to push anycast closer to optimal performance. We argue that, in the context of root DNS query resolution, such modifications and improvements would add negligible benefits to users. Specifically, we aim to show that the typical user almost never directly

queries the root DNS server. We further place an upper bound on the amount of latency a typical user of a large CDN experiences each day, and show that this latency is negligible. In light of these results, in the context of root DNS anycast deployments, any proposed improvements for anycast shown to decrease latency for users are of questionable practical utility. The rest of this paper is organized as follows. In Section 2 [a] we provide a description of the DNS infrastructure as a whole, including some practical details that are often left out of a routine presentation. In Sections 3 and 4 we attempt to provide a characterization of how root DNS latency affects the typical user, from a microscopic and macroscopic viewpoint, respectively. In Section 5 we discuss related work, and in Section 6 we conclude our findings.

2 DNS – A PRACTICAL VIEWPOINT

Modern day DNS infrastructure is designed so as to limit the time a user has to wait for DNS query resolution, and specifically root DNS query resolution.

2.1 Great Latency Savings Potential

There are many descriptions of DNS and we by no means attempt to give an introductory lesson. The DNS is a means by which user requests for human-readable hostnames are mapped to IP addresses. Typically, a user will send DNS requests in the form of UDP packets to one or more recursive resolvers (RR's) provided by their ISP¹. The RR then requests the records from a root DNS server, top level domain (TLD) server and authoritative DNS (ADNS) server corresponding to the record the user requested. Since each request is a correspondence between the RR and a remote server, there can be several requests made by the RR for a single user request.

In practice, there is quite a lot of potential for caching these DNS records, and thus removing unnecessary steps of the recursion outlined above. Each DNS server provides a time-to-live (TTL) with each record they return, specifying the duration for which the record can be kept locally cached. After this duration, the record should be deleted from the cache of the RR. It has been observed that more heavily accessed websites tend to purposefully set the TTL of their DNS records to be some small value, evidence of fine-grained traffic engineering. Hence the popularity of a website does not make it more likely to live in the cache of an RR. Nevertheless, RR's can pre-fetch records for popular websites. For example, a

configuration setting on the popular resolver BIND specifies whether or not BIND should prefetch records set to expire in a few seconds. In contrast to heavily accessed sites, TLD records tend to have long TTL's. For example, the COM TLD record (one of the most popular) has a TTL of 2 days[b][c]. One might speculate that, since websites users typically access fall into ten or so popular TLD's, that a request to the root server would rarely occur. In this scenario, a user fetching example.com would only generate a request for the COM NS record once every 2 days. Furthermore, some commercial RR's are known to cache the entire root table locally. Since there are only 1,000 or so TLD's, the memory requirements are quite inexpensive and caching these records provides huge potential for latency savings. On top of all of this structure, there can be multiple RR's (a sort of hierarchy of RR's) among which a query traverses before a third party (root, TLD, ADNS) is finally queried. This hierarchy compounds the potential effects of local caching.

Caching at remote RR's is wonderful, and can greatly reduce resolution time. Indeed since the typical user's RR is run by their local ISP, generally the user will experience negligible latency for a DNS query if they hit a warm cache. Modern browsers and operating systems have taken this one step further and engage in a number of practices that further aid the user. Operating systems or browsers are known to locally cache DNS records, as these records are small (compared to say, an image). Software can go even further, refreshing records in the cache that are about to expire (just as a RR can do). Features such as link prefetching in Chrome, where the browser will send DNS requests for all links on a page, allow the user to completely "skip" the DNS resolution process. All this being said, it is unfortunately quite difficult to quantify the extent to which these features save users time. Users request web services through an idiosyncratic combination of browsers, web sites, operating systems, personalized settings², physical locations, bandwidths, devices, and RR's. On top of this, everything that has just been mentioned can be undone through poor configuration or a conservative manager of the RR, as we will now discuss.

2.2 Practical DNS Pitfalls

We have discussed the ways in which a user is, by design, prevented from experiencing latency due to root DNS query resolution. However, there can be a stark difference between ideal and actual caching behavior. Although RR's should cache TLD records for the full

¹ The user can specify whatever RR they wish, but one can sensibly assume the typical user's RR is set by the ISP, broadcasted through DHCP.

² For example, users can opt out of the link prefetching mechanism in Chrome.

TTL, unexpected logical flows in software execution can cause unnecessary queries to the root server. Furthermore, operators can set conservative settings on the RR’s they run. Regardless of the TTL recommended by the record, the operator can set this to any value below that recommended TTL. Furthering the problem, users can disable local DNS caching on their computers, or use services such as private browsing which disables local caching for privacy concerns. Web pages can also contain “hidden” DNS requests. The HTTP body of a web page can contain references to resources stored on third party domains, for which the browser must request DNS records. Those resources can contain even more references, creating a chain of DNS requests for potentially uncached records. Users issuing reverse-DNS lookups or requesting invalid domains always result in queries to a root server (although it can be argued that these should not be considered here). Users can also act as their own RR if they wish, eliminating their ability to leverage the caching capabilities that are only possible through the “crowd-sourced” caching of typical RR’s. Again it is difficult to quantify the degree to which these pitfalls limit a typical users’ ability to leverage the caching capability of DNS. Notably it is difficult to characterize the inefficiency introduced by the software employed by the RR of that users’ ISP. One could potentially statistically characterize these inefficiencies for various open source RR’s, which is a subject for future work.

3 A CLOSE LOOK AT A RECURSIVE RESOLVER

Towards the goal of quantifying the extent to which latency due to root DNS query resolution impacts the typical user, we analyze packet traces of a recursive resolver at ISI. The recursive resolver (running BIND 9) saves all traffic traversing over port 53 to file, and has done so for 5 years, providing us with a rich source of data. This corpus (of users) is relatively small, and consists of university traffic, so the specifics of the analysis we conduct may not extend to other RR’s. Nevertheless, we wish to observe some high-level features of the data, and their implications. We analyze packet traces from the entire month of January, 2018³. We would like to quantify the effect caching root DNS records has on users, and how this differs from the ideal behavior users can experience. Specifically, we are interested in the number of queries to the root server as a fraction of user requests to the recursive resolver.

³ The same analysis for the months of May and October [d][e]resulted in the same conclusions

Assumptions	Web Page Load Time (ms)	1,000
	Root DNS Latency (ms)	500
	Number of DNS Look-Ups Per Web Page	10
Statistics	Number of Client Queries (millions)	14.9
	Number of Root Transactions	73,200
Implications	Percent of Client Queries Resulting in a Root Transaction	.49
	Expected Speed-up in PLT with No Root Latency (ms)	24.5
	Resulting PLT Speedup (percent)	2.45

Table 1: Statistics gathered from the ISI RR for the month of January, 2018, and implications of these statistics assuming conservative values for web page load time (PLT), root DNS latency, and number of DNS lookups per page.

Relevant statistics and their implications on user-perceived latency are presented in Table 1. Notably we find that, on average, a user request generates a query to the root server about .5% of the time. To obtain an upper bound on the impact of root latency, assume (conservatively) that a web page load takes 1 second, that there are 10 serial DNS requests per page, and that the root latency is 500ms⁴. The numbers used here are only exaggerated upper bounds on the impact of root DNS latency, and are only meant to provide context. Supposing root latency is completely removed from the equation a user saves, on average, 27 ms per page load. As a percentage of the total page load time, this is 2.7%, which is measurable, but not significant. To provide a sort of visual context for these results, Figure 1 shows a CDF of root DNS latency experienced for queries over the same one month period. Requests that do not generate a query to a root server are counted as having a latency of 0.

[f]

Clearly the impact of root DNS latency in this situation is minimal, and we can reasonably conclude that members of the ISI community do not experience much latency due to the root DNS servers. However, the impact on user performance is not as small as one would expect. For example, a particular day during January 2018 saw as many as 900 queries to the root server for the COM NS record. Given the 2 day TTL of this record, this query volume is absurd. This suggests that heuristic arguments that users rarely experience root latency because cached TLD records have long TTL’s are not sufficient. Manual inspection of specific query chains in

⁴ According to [http archive], and RIPE Atlas, these are quite conservative estimates.




Figure 1: Root DNS latency for queries made by clients of the ISI recursive resolver during January, 2018. Client queries that did not generate a query to a root server were given a latency of 0.

the packet traces suggests that certain logical flows in BIND can result in the root server being unnecessarily queried. We are not making claims that BIND has pathological bugs, since we did not explore the issue further. However, this is an interesting area for future work.

4 A GLOBAL LOOK AT USERS OF A LARGE CDN

Looking at cache hit rates of individual RR’s is informative, yet does not provide us with a notion of how much latency each user experiences in a day. Indeed it would be ideal if we could measure the number of active clients, how many times those clients request web pages and the cache hit rates of the clients’ recursive resolvers to obtain an estimate of the latency they experience due to root DNS resolution.

4.1 Data Sources and Processing

Motivated by this ideal scenario, we obtain statistics from a large CDN operator containing information about the users of that CDN. Specifically, these statistics include RR IP’s, the number of distinct client IP’s behind those RR’s and the relative query volume those client IP’s contributed (as a fraction of the total CDN query volume[g]) for 8 consecutive days in 2019. Working with the notion of a “user” is valuable, as this provides us with an estimate of how much latency each user experiences. However, since the notion of ‘user’ is an IP address, one ‘user’ can actually refer to several people behind a NAT.

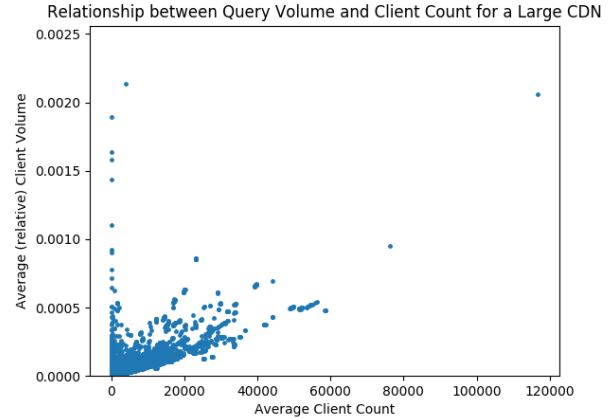


Figure 2: Relationship between daily average client count and daily average (relative) query volume for users of a large CDN. Each point represents one RR.

Furthermore, when weighting results, client query volume is intuitively the correct metric to weight by, since query volume indicates activity on the web. [h][i] The relationship between daily average query volume and daily average client IP count per RR is shown in Figure 2. The correlation coefficient between average weekly volume and client count is .81, with actual daily correlations ranging from .76 to .85. This correlation is not perfect, but good enough for us to represent the “importance” of an RR by the number of distinct client IP’s behind it.

[j] query_client_relationship

To leverage this CDN data in our aforementioned goal, we augment it with traces from the 2018 Day in the Life of the Internet (DITL), organized by OARC. DITL (run annually) contains concurrent packet captures from all root servers except for G-root⁵ for two days in Spring. From DITL, we extracted the frequency of requests each RR made to each root.

ditl_volume_comparisons

We then join these data sets by /24, aggregating their respective client IP counts, query [k]volumes, etc. . . – this joined data set is referred to henceforth as the DCDN data set. Naturally, there is a mismatch in the /24’s represented by each data set. Table 2 summarizes the extent to which the DCDN data (which is a subset of all users) represents the DITL captures, and vice versa. We also henceforth refer to these /24’s as RR’s, even though each data point may represent several RR’s. Although the

⁵ A small portion of data from the 2018 DITL is missing or anonymized. We do not compensate for this, as the impact on the results is likely small.

Impact of Root DNS Latency on User Experience

Data Set	Statistic	Percent Representation
DITL \cap CDN	DITL RR's	7.2
	DITL Volume	45.9
	CDN RR's	87.7
	CDN Volume	88.9
DITL \cap CDN \cap RIPE	DITL RR's	.1
	DITL Volume	14.4
	CDN RR's	.9
	CDN Volume	53.9

Table 2: Statistics displaying the extent to which the RR's of users in a large CDN represent RR's seen in the 2018 DITL captures. Further, it shows the extent to which RR's of RIPE probes represent the 2018 DITL captures.

DCDN data considers a relatively small percentage of all RR's, it captures a disproportionately large amount all DITL volume. A useful visualization of the effect of removing these RR's from consideration is shown in Figure ?? . The quartiles of the DCDN volume counts (blue) are approximately 10-times those of DITL (orange), suggesting /24s in the DCDN particularly active subset of all RR's. This suggests that, despite disregarding many RR's, our analysis actually provides an upper bound on root DNS latency experienced by the typical user.

4.2 Root Latency Experienced by Users

Our main result, shown in Figure 3, is a CDF of expected user latency per day, where the expected value is calculated according to certain assumptions we make about root latency. We have also included the number of requests each user makes per day in Figure ?? for reference. Figure ?? may be particularly useful for comparing these results with other studies that show CDF's of user latencies to various roots.[l][m]

user_root_latency_per_day

To generate each line in Figure 3, the expected root latency per RR per query is multiplied by the number of queries per day each RR makes. This is then weighted by user count (from the CDN data set) and the resulting CDF is calculated. The only question left to answer is how to calculate the expected root latency per RR. As a useful comparison, we have included lines labeled “basic-n” which are generated naively assuming the latency from every RR to every root server is n ms. Similarly, the line labeled “basic-80-everything” is generated further assuming every RR in DITL, but not in the CDN data set, has one user per day. This line therefore includes information about every RR seen in DITL and is a complete overestimation of actual latency experienced




Figure 3: A CDF of approximate latency a user experiences due to root DNS resolution, per day.

by each user; however, the two lines are barely distinguishable. This is further proof of the representativeness of the DCDN data set.

The line labeled “heuristic root value” are calculated as follows: assume each RR queries the set of roots according to a distribution p_{RR} , where p_{RR} is a discrete p.d.f. over the 12 root servers indicating the probability with which the RR chooses each root server to query. This p.d.f. is calculated directly from the DITL data. Next, obtain median latency values to each root (e.g. from RIPE Atlas) and call these l_{RR} . Note that, for now, l_{RR} is independent of RR. Compute the expected latency from each RR to any root as the dot product between l_{RR} and p_{RR} (i.e. the expectation of the latency). Although somewhat imprecise, this method of calculating latency for each user takes into account that some root servers generally exhibit lower latencies than others due to investments in infrastructure⁶. Hence, this line can be considered a more reasonable estimate of the latency users experience due to root DNS resolution. This is corroborated by the “middle-of-the-pack” median and 95th percentile estimates of .352ms/day and 4.104ms/day, respectively. Finally, the line labeled “RIPE” looks only at those RR's in both the DITL and CDN data sets housing RIPE probes. RIPE probes routinely issue ping measurements to each root server and report latency values. From the preceding paragraph, l_{RR} is then populated from these measurements for each RR, and the expected root latency per query is calculated per RR according to p_{RR} . From Table 2, we see that /24's housing RR's of RIPE

⁶ For example, at the time of writing the median latency to B root (3 sites) is 130ms whereas the median latency to F root (225 sites) is less than 5 ms.

probes only account for 14.4% of the volume in DITL, and less than .1% of all /24's; hence, RIPE probes are not representative of the population as expected. However, for clients in these /24's, this is a fairly accurate measurement of the expected daily latency due to root DNS resolution. We see this method provides the lowest median and 95th percentile estimates of .225ms/day and 4.011ms/day, respectively. This makes sense, as RIPE probes probably reside in areas of the world with better connectivity.

Regardless of which method is used to calculate expected latency, users generally experience median values of < 1 ms/day while worst case users generally experience < 50 ms/day. Even when making the conservative assumption that every root query from every RR takes 500ms, 95% of users experience no more than 80ms/day of latency due to root DNS resolution which is *inconsequential*. Moreover, recall that the notion of “user” here is likely a NAT, representing several users, and that we do not prune requests unrelated to web traffic (e.g. PTR requests) or invalid domain look-ups when parsing DITL traces. This further establishes that Figure 3 is to be interpreted as a conservative upper bound of daily root latency experienced by users. Although it may not be valid to reason about specific numbers obtained from the above analysis, we can conservatively say that the typical user hardly ever experiences latency due to root DNS resolution. Therefore when assessing the utility or performance of IP anycast, it may not be sufficient or useful to point at a set of root DNS queries that incurred unnecessarily high latency. As we show here, users rarely interact with the root DNS infrastructure. Hence the utility of proposed improvements to IP anycast should not be assessed using root DNS servers, as the generalization ability of the improvement to all IP anycast systems is unpredictable.

5 RELATED WORK[N]

IP anycast is usually studied in two applications: the root DNS servers [1], and CDN's [2]. In addition to these topics, we discuss studies of popular recursive resolvers, and client-centric measurements of web latency.

5.1 Root DNS Anycast Performance

The performance of anycast in the context of root DNS is generally gauged by anycast's ability to balance load among server replicas or provide low latency to users. Generally, all studies conclude that anycast successfully balances load, while latency performance depends on the specific deployment configuration. [14] looks at a DDoS

attack on the root name server infrastructure, and generally shows that anycast is a good defense mechanism against such attacks. An earlier study, [15] confirms that anycast protects the root DNS infrastructure against such attacks and, furthermore, that anycast routes users to an optimal location in most cases. [7] looks at user latency to C, F, K, and L-root and attributes better performance to good geographic location and peering strategies. These findings coincide with an earlier study, [3], who conclude the performance of anycast is intrinsically linked to deployment strategy. Additionally [7] finds that as few as 12 sites can provide “good” latency to users. [1], [6], [7], and [13] are all examples of studies who quantify latencies to various root servers, and note how these compare to the (optimal) latency of the closest unicast alternative for the client who issued the query.

5.2 CDN Anycast Performance

Some CDN's (e.g. Cloudflare, Edgecast, Fastly) use IP anycast to augment their serving infrastructure. When deploying an Anycast CDN (ACDN), delivering content to users with low latency becomes a high priority, as there is a large financial incentive to do so. The simplicity of IP anycast comes at the cost of having coarse grained control over where client queries land. Shifting client load between nodes during peak hours, for example, is a challenging problem. As a potential solution, [9] and [1] use DNS redirects at ADNS servers to shift load among anycast nodes, albeit in slightly different ways. [4] analyzes what latency clients are achieving, compared to optimal, when being routed to anycast nodes and finds that 10% of clients experience a latency inflation of at least 100 ms.

5.3 Recursive Resolvers and the Benefits of Caching

Similar to the RR analysis conducted here, [10] looks at DNS traffic on a small network and notably finds that 16% of queries resulted in queries to the root, most of which were for invalid domains. As this study is quite old, it is no surprise that this rate has decreased (recall we observed .5% of queries resulted in queries to the root) since browser designers and network engineers understand the importance of caching. [5] also look at a RR network and analyze statistics of DNS exchanges occurring over it including DNS transaction latencies. Both [11] and [17] look at certain pathological behaviors of popular recursive resolvers, and the implications these behaviors have on root DNS load.

5.4 Web Performance

Although we were unable to find any specific study that looked at how web performance and root DNS latency were related, there are certainly studies characterizing web performance. [16] characterizes web performance bottlenecks in (at the time) new broadband networks, and finds that latency is the main bottleneck for PLT when the user’s bandwidth exceeds 16 Mbps. However, the study does not realistically emulate a page load and, in particular, can not analyze the effect of having multiple DNS resolutions per page. Similarly, [2] analyzes how each step of a page load contributes to the aggregate PLT using a tool designed in-house. However, unlike [16], they did not conduct a large measurement campaign and do not include information about multiple DNS lookups per page. A more recent study, [8] provides a brief survey of web performance measurement studies and explains why it is difficult (with current practices) to compare two different studies in web performance.

6 CONCLUSION

IP anycast has come under attack, with studies showing, for example, how BGP can naturally route clients to suboptimal anycast instances and inflate client latencies. Due to the relative availability of root DNS data and diverse deployment strategies of the root DNS servers, they are common targets for delineating inefficiencies and suggesting improvements to IP anycast. We argue not only that the root DNS servers have different design goals (i.e. resiliency against attacks) than that of other anycast services, but also that users rarely interact with the root DNS infrastructure – rendering perceived inefficiencies and proposed improvements to be ill-founded when only tested on the root DNS. Perhaps simple yet effective ideas such as browser link prefetching or DNS request parallelization should be expanded and their adoption by users encouraged, rather than proposed improvements to IP anycast. [a]section labels is a better way to go about it maybe [b]maybe generate a figure showing a CDF of TTL’s of all TLD NS records (pretty easy to generate) [c]Would probably need to run BIND locally, and set myself as my own recursive resolver [d]has 4x cache hit rate (far fewer client queries?) [e]Update: pinged John about it with more information about when issue occurs [f]maybe use this as a p.d.f. for “amortizing” the root latency in the table [g]perhaps elaborate on what is meant by “query volume” [h]+ethanbkb@gmail.com This is one way I currently use Microsoft’s query volume numbers to justify our results. What do you think? [i]I was thinking along the lines of weighting root query cost by query volume (rather than by “users”). [j]y axis needs

relabeling [k]TODO: comment about removing v6 [l]For example, people could look at UMD paper and apply the 95%-ile latency to all clients and see implications. [m]Questionable utility, may be a glut of graphs [n]cite more numbers maybe [o]interesting that in 10 minutes they observe so many queries for COM TLD yet don’t see any issue with that [p]Might be an interesting tool to use

REFERENCES

- [1] H. A. Alzoubi, S. Lee, M. Rabinovich, O. Spatscheck, and J. Van Der Merwe. 2011. A practical architecture for an anycast CDN. *ACM Transactions on the Web (TWEB)* 5, 4 (2011), 17.
- [2] A. S. Asrese, P. Sarolahti, M. Boye, and J. Ott. WePR: a tool for automated web performance measurement. In *2016 IEEE Globecom Workshops (GC Wkshps)*. IEEE, 1–6 2016.
- [3] H. Ballani, P. Francis, and S. Ratnasamy. A measurement-based deployment proposal for IP anycast. In *Proceedings of the 6th ACM SIGCOMM conference on Internet measurement.*, ACM, 231–244 2006.
- [4] M. Calder, A. Flavel, E. Katz-Bassett, R. Mahajan, and J. Padhye. Analyzing the Performance of an Anycast CDN. In *Proceedings of the 2015 Internet Measurement Conference.*, ACM, 531–537 2015.
- [5] T. Callahan, M. Allman, and M. Rabinovich. 2013. On modern DNS behavior and properties. *ACM SIGCOMM Computer Communication Review* 43, 3 (2013), 7–15.
- [6] L. Colitti, E. Romijn, H. Uijterwaal, and A. Robachevsky. 2006. Evaluating the effects of anycast on DNS root name servers. *RIPE document RIPE-393* 6 (2006).
- [7] R. de Oliveira Schmidt, J. Heidemann, and J. H. Kuipers. Anycast latency: How many sites are enough?. In *International Conference on Passive and Active Network Measurement.*, Springer, 188–200 2017.
- [8] T. Enghardt, T. Zinner, and A. Feldmann. Web performance pitfalls. In *International Conference on Passive and Active Network Measurement.*, Springer, 286–303 2019.
- [9] A. Flavel, P. Mani, D. Maltz, N. Holt, J. Liu, Y. Chen, and O. Surmachev. Fastroute: A scalable load-aware anycast routing architecture for modern cdns. In *12th {USENIX} Symposium on Networked Systems Design and Implementation ({NSDI} 15)*. 381–394 2015.
- [10] J. Jung, E. Sit, H. Balakrishnan, and R. Morris. 2002. DNS performance and the effectiveness of caching. *IEEE/ACM Transactions on networking* 10, 5 (2002), 589–603.
- [11] M. Lentz, D. Levin, J. Castonguay, N. Spring, and B. Bhattacharjee. D-mystifying the D-root Address Change. In *Proceedings of the 2013 conference on Internet measurement conference.*, ACM, 57–62 2013.
- [12] Z. Li, D. Levin, N. Spring, and B. Bhattacharjee. Internet anycast: performance, problems, & potential.. In *SIGCOMM.*, 59–73 2018.
- [13] J. Liang, J. Jiang, H. Duan, K. Li, and J. Wu. Measuring query latency of top level DNS servers. In *International Conference on Passive and Active Network Measurement.*, Springer, 145–154 2013.
- [14] G. Moura, R. d. O. Schmidt, J. Heidemann, W. B. de Vries, M. Muller, L. Wei, and C. Hesselman. Anycast vs. DDoS: Evaluating the November 2015 root DNS event. In *Proceedings*

John Heidemann (ISI), Matt Calder (Microsoft Research), Arpit Gupta (UCSB), Ethan Katz-Bassett (Columbia University), and Thomas Koch (Columbia University)

- of the 2016 Internet Measurement Conference*,. ACM, 255–270 2016.
- [15] S. Sarat, V. Pappas, and A. Terzis. On the use of anycast in DNS. In *Proceedings of 15th International Conference on Computer Communications and Networks*,. IEEE, 71–78 2006.
- [16] S. Sundaresan, N. Magharei, N. Feamster, R. Teixeira, and S. Crawford. 2013. Web performance bottlenecks in broadband access networks. *ACM SIGMETRICS Performance Evaluation Review* 41, 1 (2013), 383–384.
- [17] Y. Yu, D. Wessels, M. Larson, and L. Zhang. 2012. Authority server selection in DNS caching resolvers. *ACM SIGCOMM Computer Communication Review* 42, 2 (2012), 80–86.